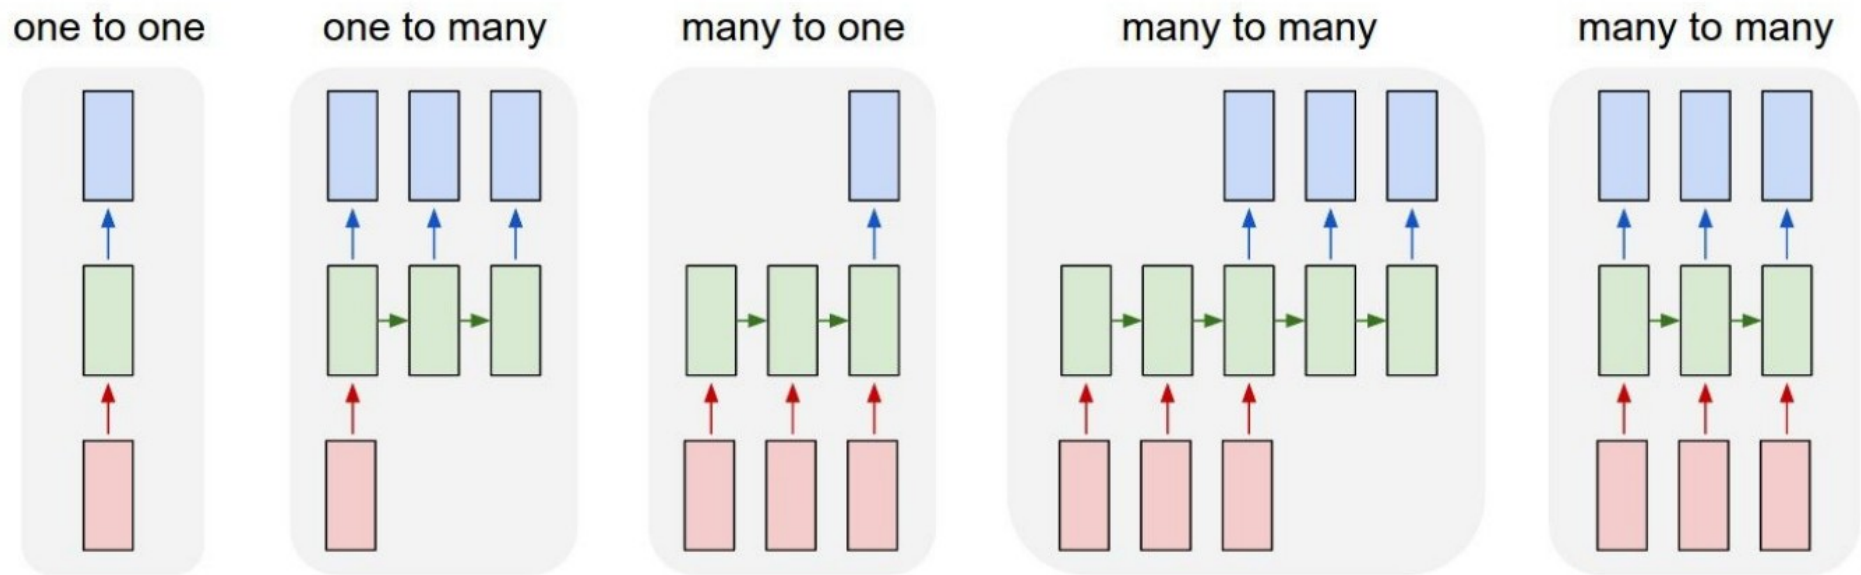


Machine Learning

- Intro to Recurrent Neural Networks -

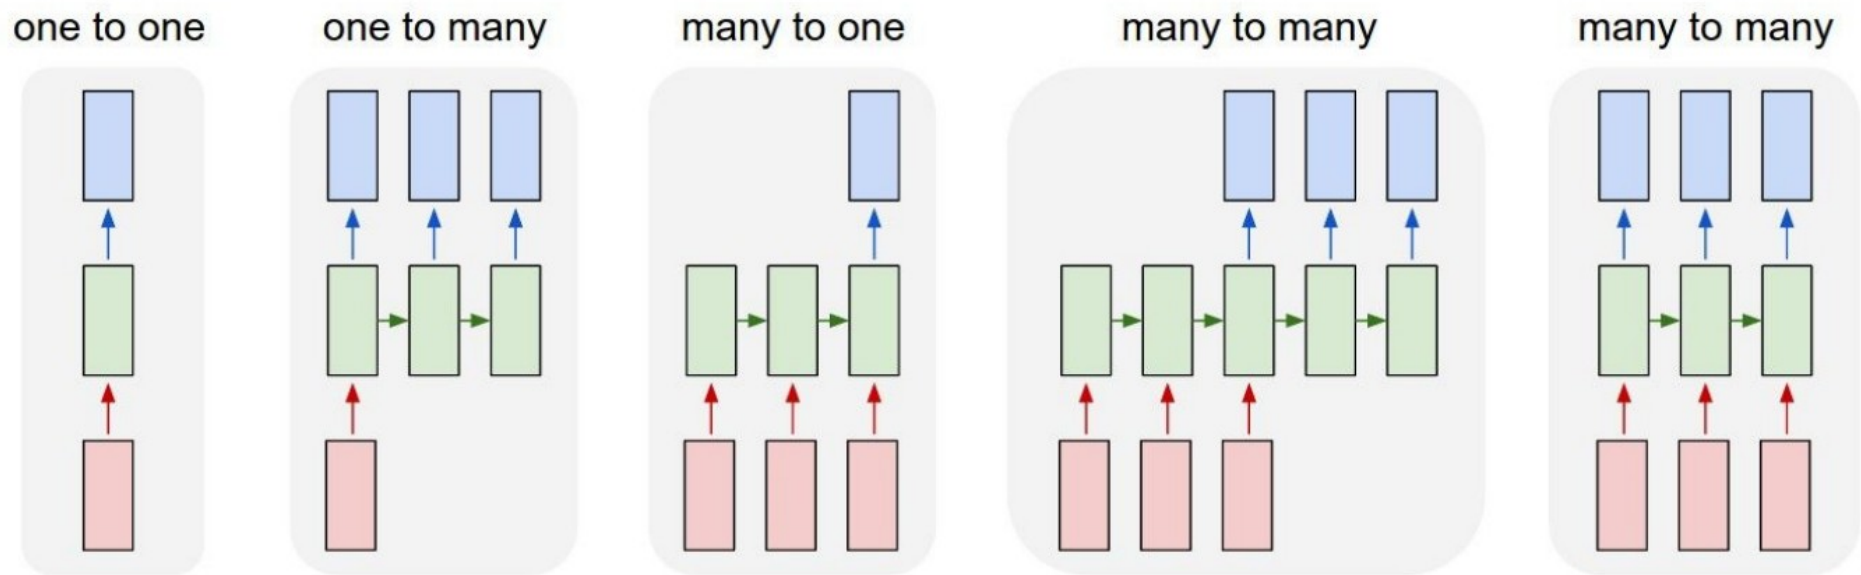
RNN Tasks

Recurrent Neural Networks: Process Sequences



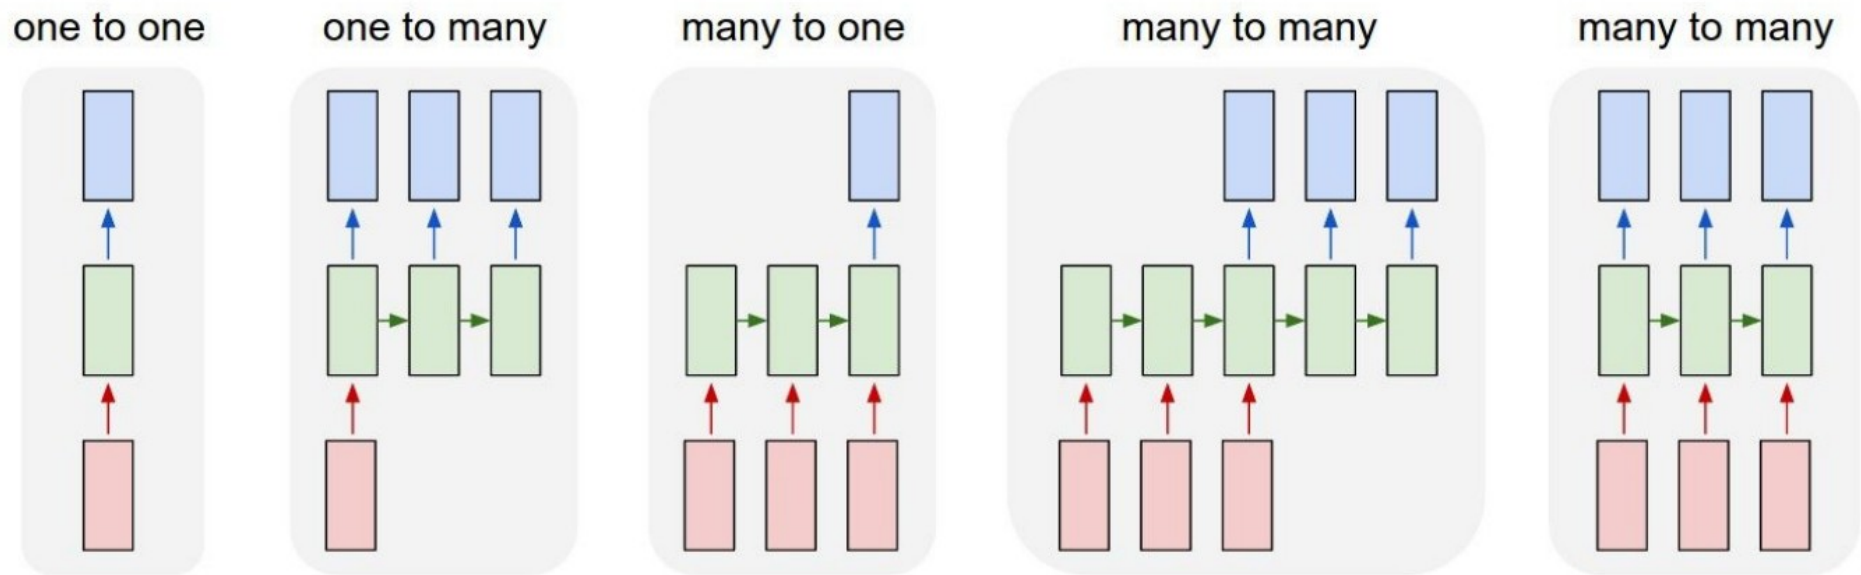
Vanilla RNNs

Recurrent Neural Networks: Process Sequences



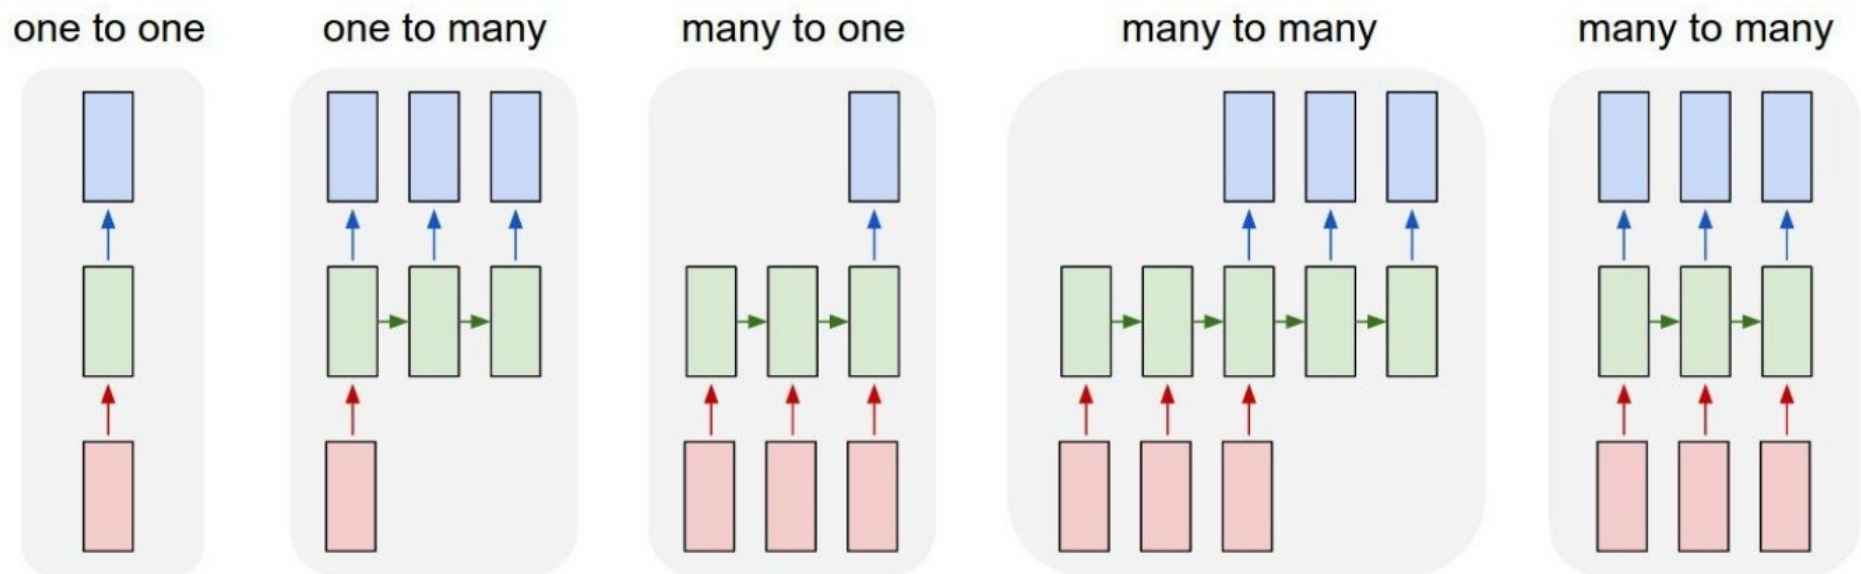
e.g. Image Captioning
Image → sequence of words

Recurrent Neural Networks: Process Sequences



e.g. Sentiment Classification
Sequence of words → sentiment

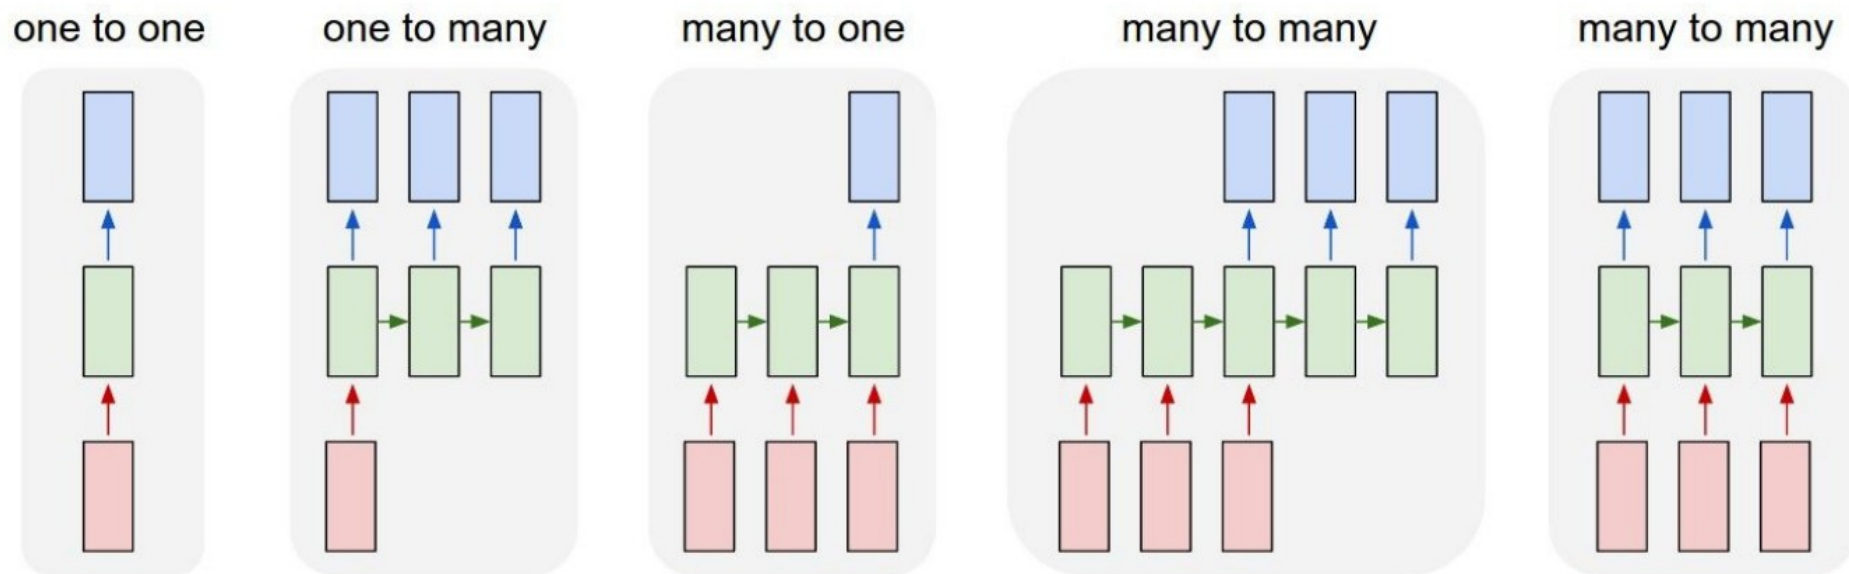
Recurrent Neural Networks: Process Sequences



e.g. Translation

Sequence of words → sequence of words

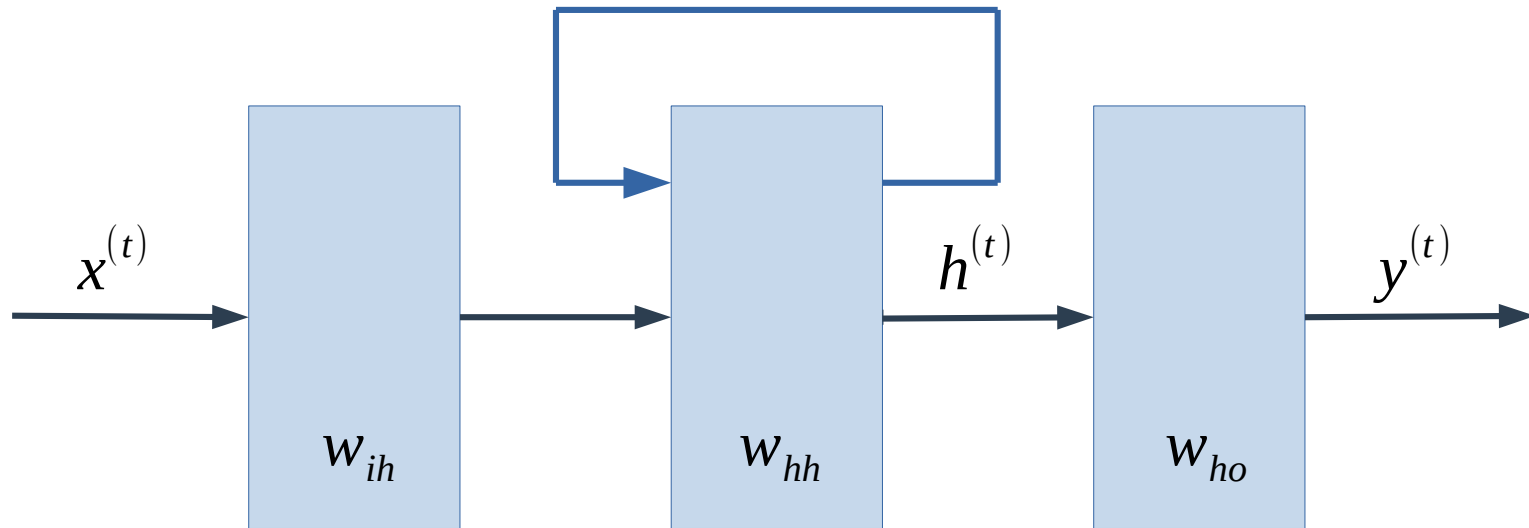
Recurrent Neural Networks: Process Sequences



e.g. Video classification
on frame level

RNN Model

Vanilla RNN Model

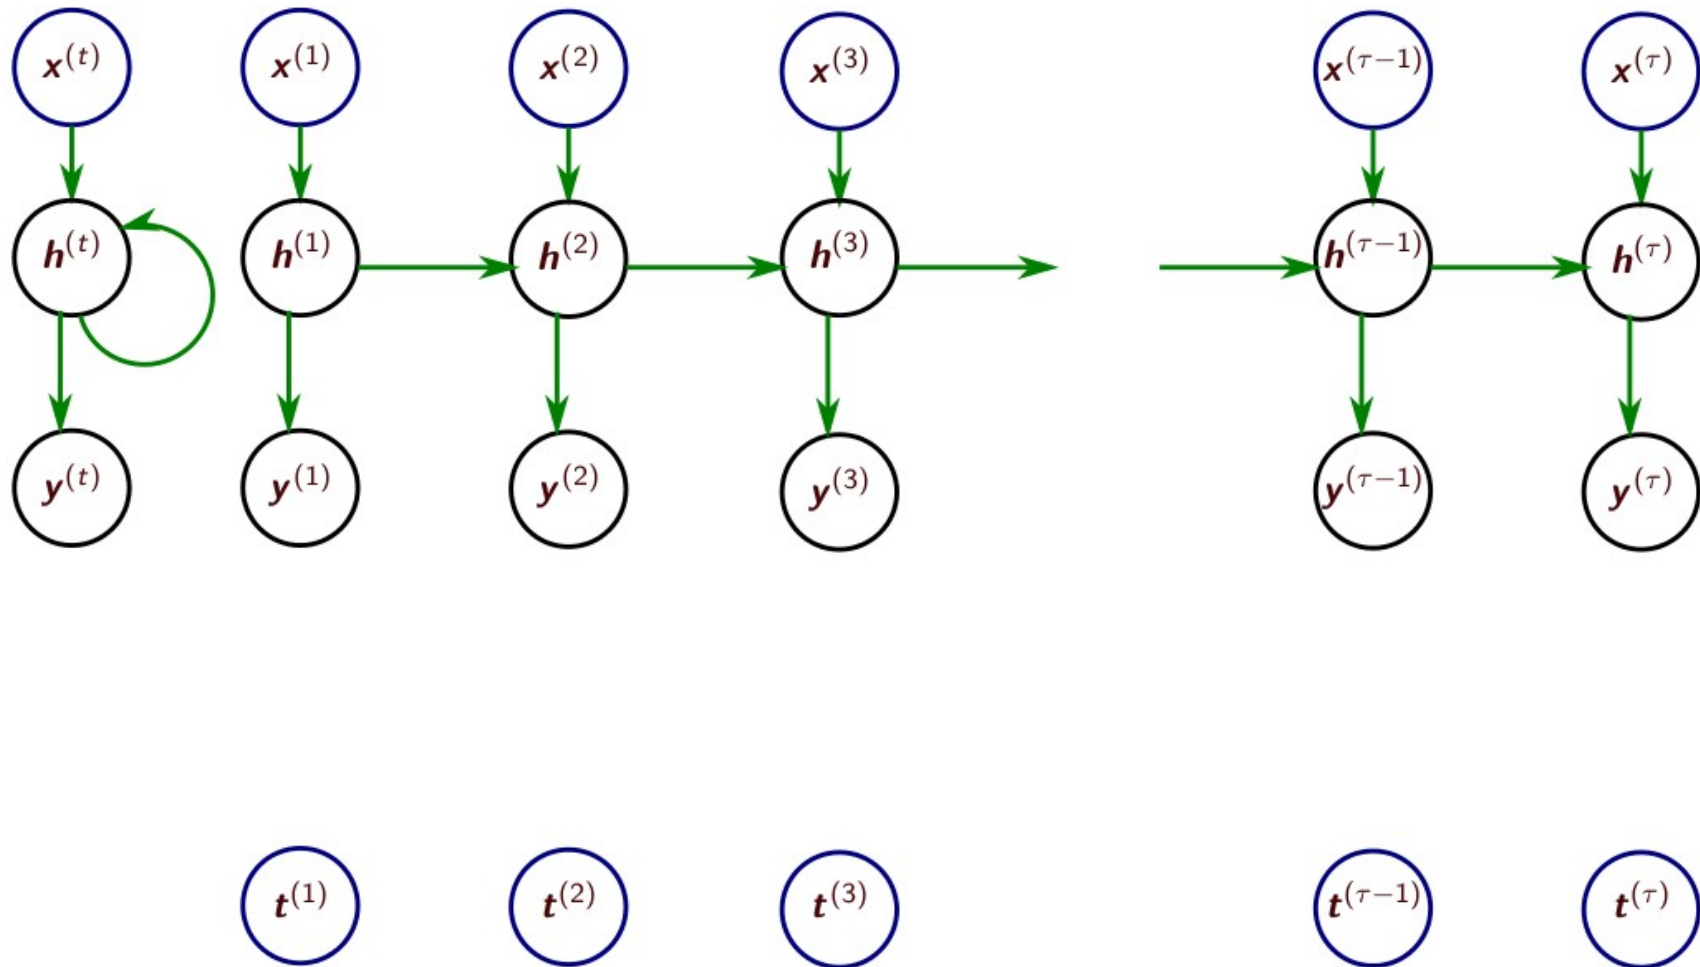


- Current state depends on *current inputs* and *previous state*
- RNNs can yield outputs at each time step

$$h^{(t)} = f_{w_{hh}}(h^{(t-1)}, f_{w_{ih}}(x^{(t)}))$$

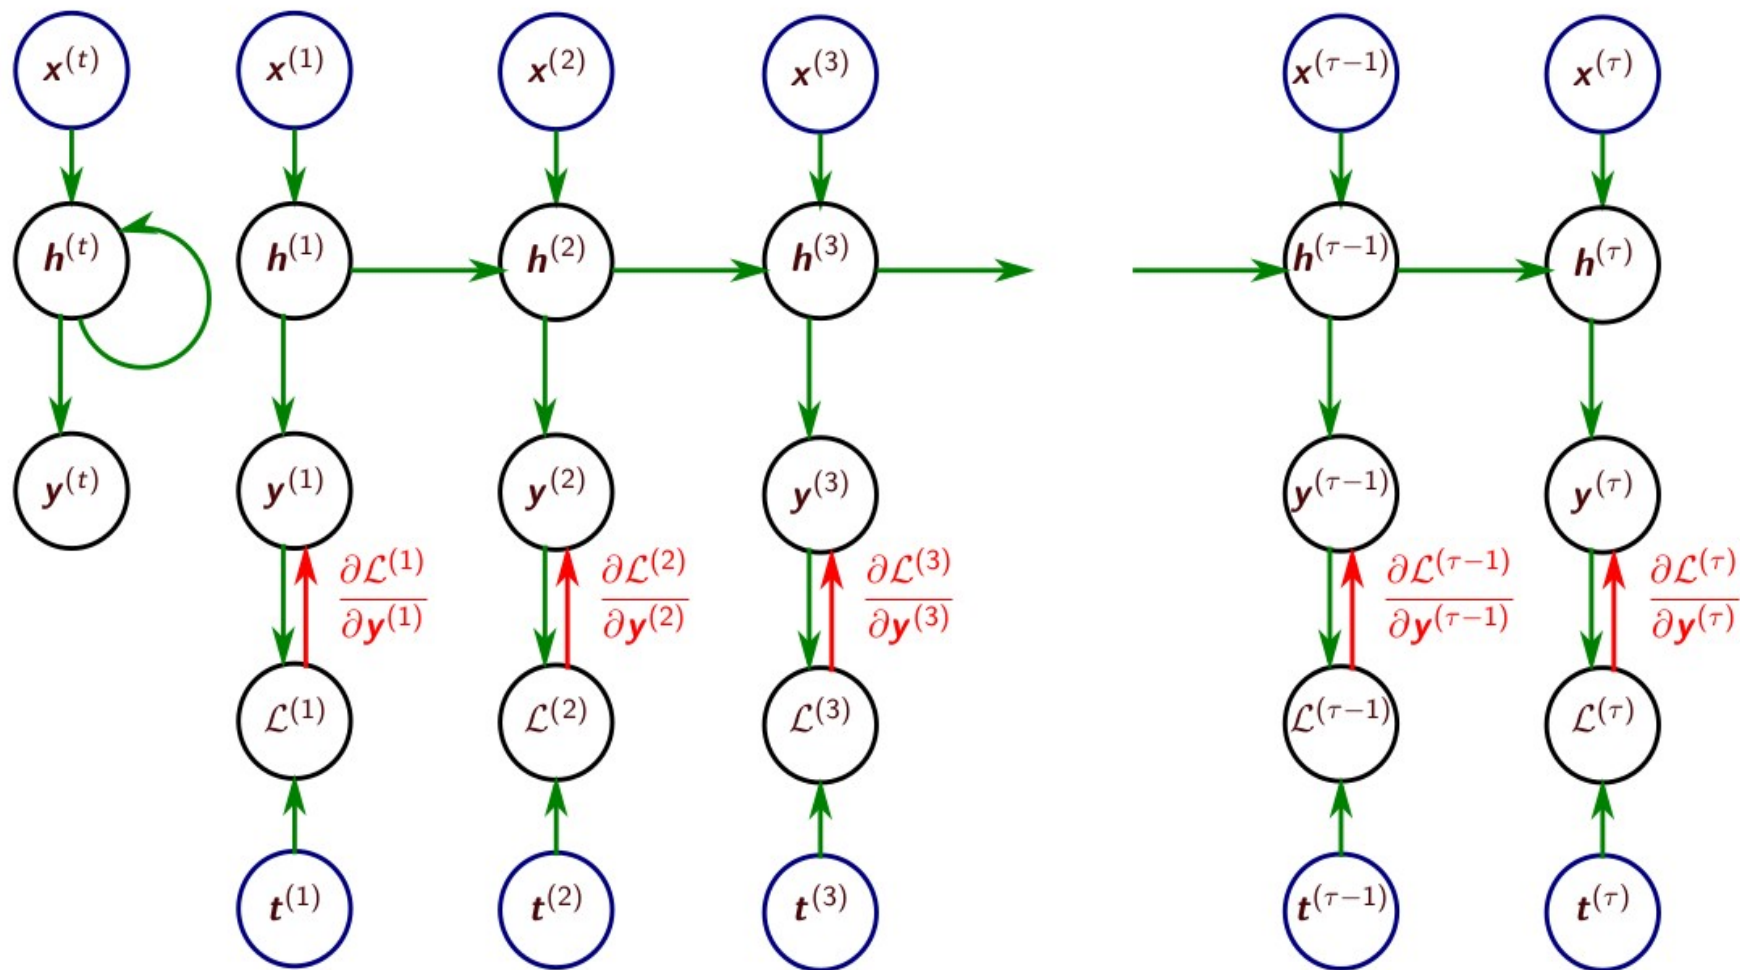
$$y^{(t)} = f_{w_{ho}}(h^{(t)}), \forall t \in \{1 \dots \tau\}$$

Unfolding RNN in time



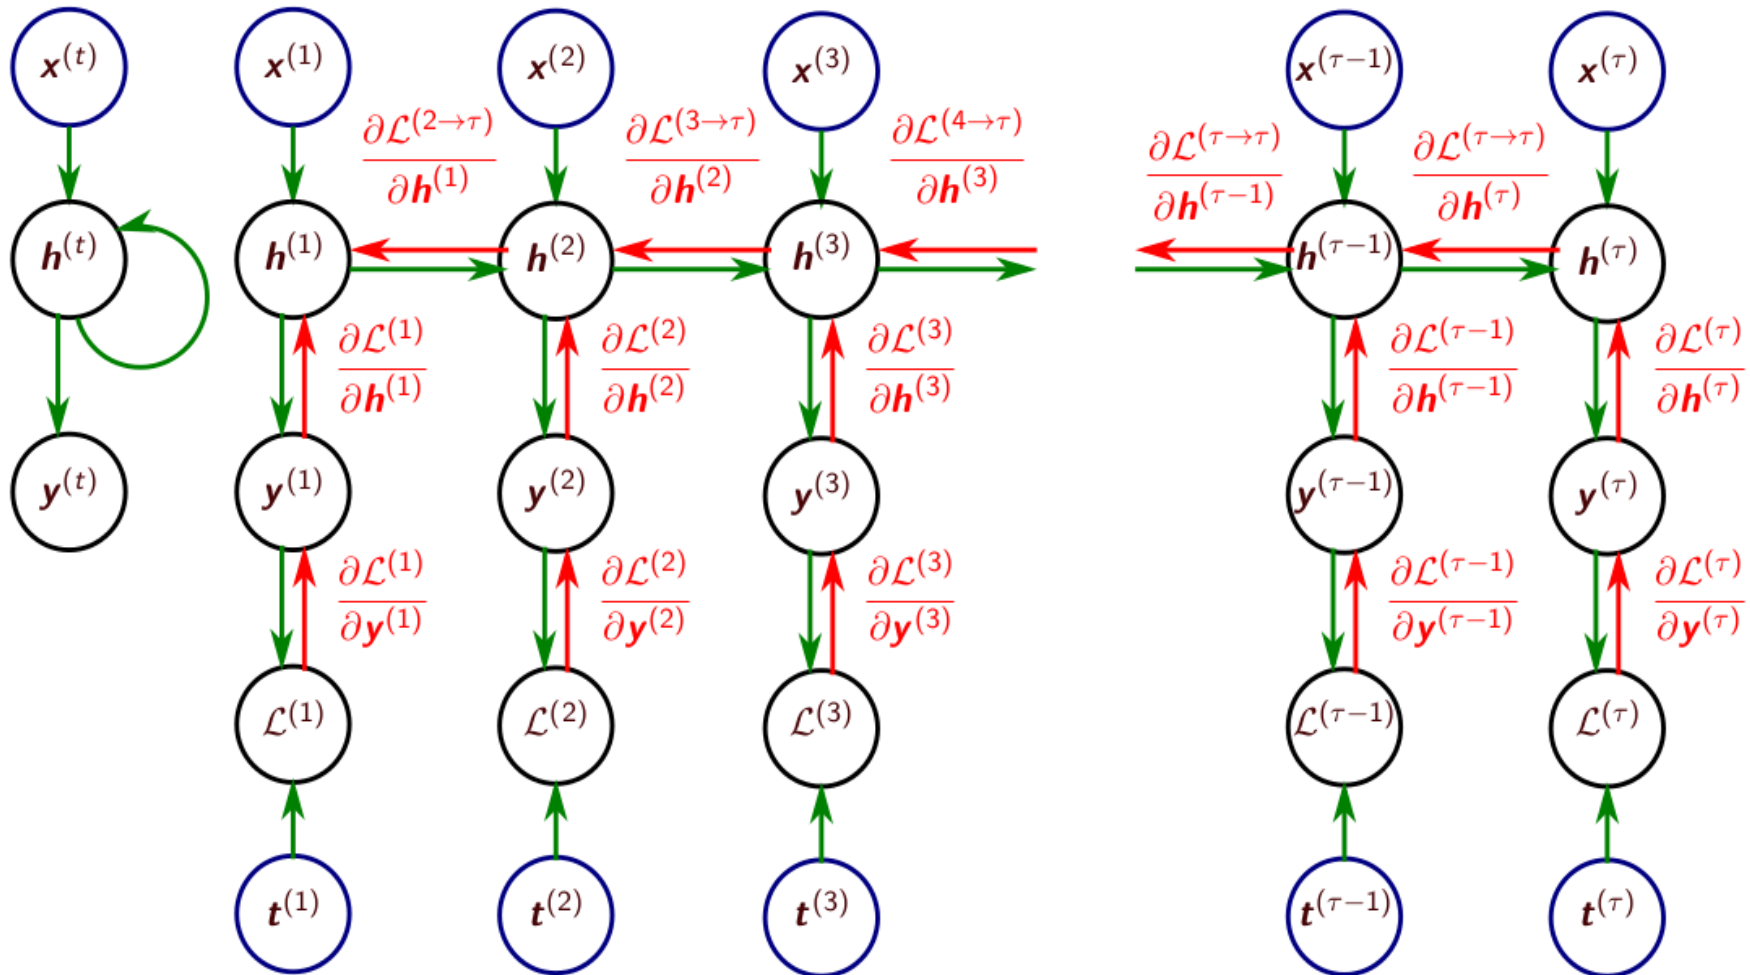
Source: NN Lectures, Tudor
Berariu, 2016

Unfolding RNN in time



Source: NN Lectures, Tudor Berariu, 2016

Unfolding RNN in time



Source: NN Lectures, Tudor Berariu, 2016

Truncated BPTT

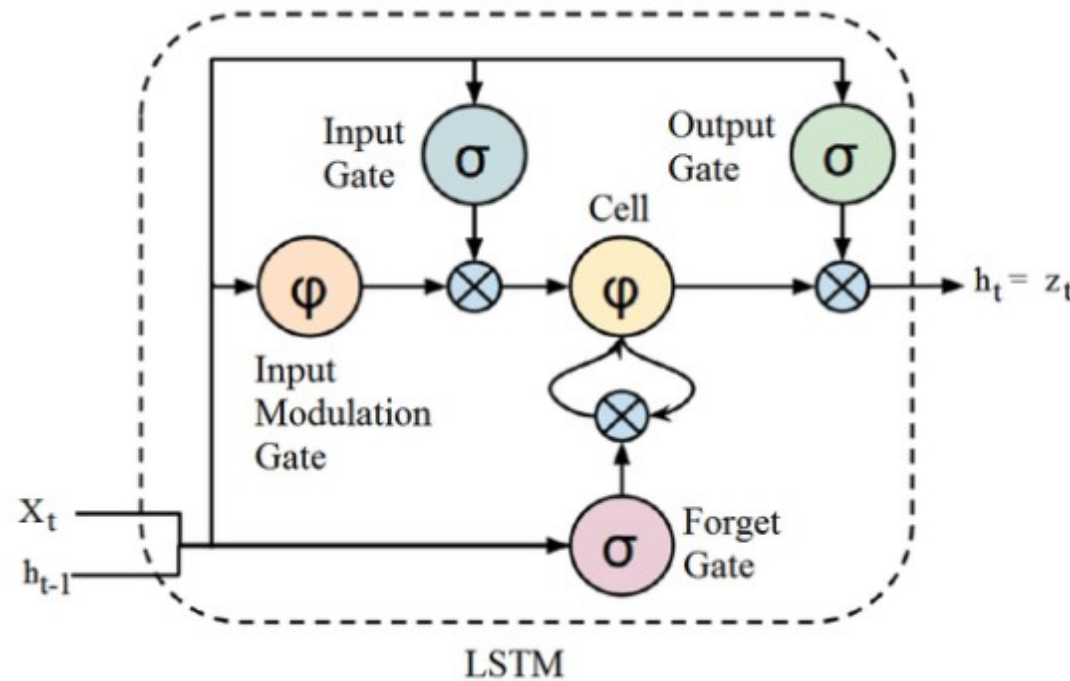
- Used in practice
- Summary of the algorithm:
 - Present a sequence of k_1 timesteps of input and output pairs to the network.
 - Unroll the network then calculate and accumulate errors across k_2 timesteps.
 - Roll-up the network and update weights.
 - Repeat

Teacher Forcing and Warm-start

- When training a RNN to generate a sequence, often, the predictions (outputs $y^{(t)}$) of a RNN cell are used as the input of the cell at the next timestamp
- **Teacher Forcing**: at training time, use the ***targets*** of the sequence, instead of RNN predictions, as inputs to the next step
- **Warm-start**: when using an RNN to predict a next value conditioned on *previous* predictions, it is sometimes necessary to give the RNN some *context* (known ground truth elements) before letting it predict on its own

LSTM

LSTM Cell



Img source:
<https://medium.com/@kangeengine/>

- Input Gate (***i*** in $(0, 1)$ – sigmoid) – scales input to cell (write)
- Output Gate (***o*** in $(0, 1)$ – sigmoid) – scales output from cell (read)
- Forget Gate (***f*** in $(0, 1)$ – sigmoid) – scales old cell values (reset mem)

LSTM Cell - Equations

$$i_t = \sigma(\theta_{xi} x^{(t)} + \theta_{hi} h^{(t-1)} + b_i)$$

$$f_t = \sigma(\theta_{xf} x^{(t)} + \theta_{hf} h^{(t-1)} + b_f)$$


$$o_t = \sigma(\theta_{xo} x^{(t)} + \theta_{ho} h^{(t-1)} + b_o)$$

$$g_t = \tanh(\theta_{xg} x^{(t)} + \theta_{hg} h^{(t-1)} + b_g)$$

$$c_t = f_t \odot c_{(t-1)} + i_t \odot g_t$$

$$h_t = o_t \odot \tanh(c_t), \text{ where } \odot \text{ is elementwise multiplication}$$

LSTMs in practice



- Sutskever et al, Sequence to Sequence Learning with Neural Networks, NIPS 2014
 - Models are huge :-)
 - 4 layers, 1000 LSTM cells per layer
 - Input vocabulary of 160k
 - Output vocabulary of 80k
 - 1000 dimensional word embeddings